

Team ID: PNT2022TMID07706

FERTILIZER RECOMMENDATION SYSTEM FOR DISEASE PREDECTION.

Crop recommendation model

```
In [1]: # Importing libraries

from __future__ import print_function
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import classification_report
from sklearn import metrics
from sklearn import tree
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: df = pd.read_csv('../Data-processed/crop-recommendation.csv')
```

```
In [3]: df.head()
```

```
Out[3]:
```

	N	P	K	temperature	humidity	ph	rainfall	label
0	90	42	43	20.879744	82.002744	6.502985	202.935536	rice
1	85	58	41	21.770462	80.319644	7.038096	226.655537	rice
2	60	55	44	23.004459	82.320763	7.840207	263.964248	rice
3	74	35	40	26.491096	80.158363	6.980401	242.864034	rice
4	78	42	42	20.130175	81.604873	7.628473	262.717340	rice

```
In [4]: df.tail()
```

```
Out[4]:
```

	N	P	K	temperature	humidity	ph	rainfall	label
2195	107	34	32	26.774637	66.413269	6.780064	177.774507	coffee

2196	99	15	27	27.417112	56.636362	6.086922	127.924610	coffee
2197	118	33	30	24.131797	67.225123	6.362608	173.322839	coffee
2198	117	32	34	26.272418	52.127394	6.758793	127.175293	coffee
2199	104	18	30	23.603016	60.396475	6.779833	140.937041	coffee

```
In [5]: df.size
```

```
Out[5]: 17600
```

```
In [6]: df.shape
```

```
Out[6]: (2200, 8)
```

```
In [7]: df.columns
```

```
Out[7]: Index(['N', 'P', 'K', 'temperature', 'humidity', 'ph', 'rainfall', 'label'], dtype='object')
```

```
In [8]: df['label'].unique()
```

```
Out[8]: array(['rice', 'maize', 'chickpea', 'kidneybeans', 'pigeonpeas',
       'mothbeans', 'mungbean', 'blackgram', 'lentil', 'pomegranate',
       'banana', 'mango', 'grapes', 'watermelon', 'muskmelon', 'apple',
       'orange', 'papaya', 'coconut', 'cotton', 'jute', 'coffee'],
      dtype=object)
```

```
In [9]: df.dtypes
```

```
Out[9]: N          int64
P          int64
K          int64
temperature  float64
```

```

temperature    float64
humidity       float64
ph             float64
rainfall       float64
label          object
dtype: object

```

```
In [10]: df['label'].value_counts()
```

```


Out[10]: muskmelon    100
kidneybeans    100
papaya        100
pigeonpeas    100
blackgram     100
cotton        100
mothbeans     100
mungbean      100
watermelon    100
orange        100
mango        100
banana        100
rice          100
pomegranate   100
chickpea      100
apple         100
jute          100
grapes        100
lentil        100
coffee       100
maize         100
coconut       100
Name: label, dtype: int64

```

```
In [11]: sns.heatmap(df.corr(),annot=True)
```

```
Out[11]:
```

```
Out[11]:
```



Seperating features and target label

```

In [12]: features = df[['N', 'P','K','temperature', 'humidity', 'ph', 'rainfall']]
target = df['label']
#features = df[['temperature', 'humidity', 'ph', 'rainfall']]
labels = df['label']

```

```

In [13]: # Initializing empty lists to append all model's name and corresponding name
acc = []
model = []

```

```
In [14]: # Splitting into train and test data

from sklearn.model_selection import train_test_split
Xtrain, Xtest, Ytrain, Ytest = train_test_split(features, target, test_size = 0.2, random_state = 2)
```

Decision Tree

```
In [15]: from sklearn.tree import DecisionTreeClassifier

DecisionTree = DecisionTreeClassifier(criterion="entropy", random_state=2, max_depth=5)

DecisionTree.fit(Xtrain, Ytrain)

predicted_values = DecisionTree.predict(Xtest)
x = metrics.accuracy_score(Ytest, predicted_values)
acc.append(x)
model.append('Decision Tree')
print("DecisionTrees's Accuracy is: ", x*100)

print(classification_report(Ytest, predicted_values))
```

```
DecisionTrees's Accuracy is: 90.0
precision    recall  f1-score   support

   apple      1.00      1.00      1.00        13
  banana      1.00      1.00      1.00        17
 blackgram    0.59      1.00      0.74        16
  chickpea     1.00      1.00      1.00        21
   coconut    0.91      1.00      0.95        21
   coffee     1.00      1.00      1.00        22
   cotton     1.00      1.00      1.00        20
   grapes     1.00      1.00      1.00        18
    jute       0.74      0.93      0.83        28
 kidneybeans  0.00      0.00      0.00        14
   lentil     0.68      1.00      0.81        23

   maize      1.00      1.00      1.00        21
   mango      1.00      1.00      1.00        26
  mothbeans   0.00      0.00      0.00        19
  mungbean     1.00      1.00      1.00        24
 muskmelon    1.00      1.00      1.00        23
   orange     1.00      1.00      1.00        29
   papaya     1.00      0.84      0.91        19
 pigeonpeas   0.62      1.00      0.77        18
 pomegranate  1.00      1.00      1.00        17
    rice      1.00      0.62      0.77        16
 watermelon   1.00      1.00      1.00        15

 accuracy          0.90      440
 macro avg         0.84      0.88      0.85      440
 weighted avg       0.86      0.90      0.87      440
```

```
In [16]: from sklearn.model_selection import cross_val_score
```

```
In [17]: # Cross validation score (Decision Tree)
score = cross_val_score(DecisionTree, features, target, cv=5)
```

```
In [18]: score
```

```
Out[18]: array([0.93636364, 0.90909091, 0.91818182, 0.87045455, 0.93636364])
```

Saving trained Decision Tree model

```
In [19]: import pickle
# Dump the trained Naive Bayes classifier with Pickle
DT_pkl_filename = '../models/DecisionTree.pkl'
# Open the file to save as pkl file
DT_Model_pkl = open(DT_pkl_filename, 'wb')
pickle.dump(DecisionTree, DT_Model_pkl)
# Close the pickle instances
DT_Model_pkl.close()
```

Guassian Naive Bayes

```
In [20]: from sklearn.naive_bayes import GaussianNB

NaiveBayes = GaussianNB()

NaiveBayes.fit(Xtrain,Ytrain)

predicted_values = NaiveBayes.predict(Xtest)
x = metrics.accuracy_score(Ytest, predicted_values)
acc.append(x)
model.append('Naive Bayes')
print("Naive Bayes's Accuracy is: ", x)

print(classification_report(Ytest,predicted_values))
```

```
Naive Bayes's Accuracy is: 0.99090909090909091
precision    recall  f1-score   support

   apple      1.00      1.00      1.00        13
  banana      1.00      1.00      1.00        17
 blackgram      1.00      1.00      1.00        16
```

```
   coffee      1.00      1.00      1.00        22
   cotton      1.00      1.00      1.00        20
   grapes      1.00      1.00      1.00        18
    jute       0.88      1.00      0.93        28
 kidneybeans    1.00      1.00      1.00        14
   lentil      1.00      1.00      1.00        23
   maize       1.00      1.00      1.00        21
   mango       1.00      1.00      1.00        26
 mothbeans     1.00      1.00      1.00        19
 mungbean      1.00      1.00      1.00        24
 muskmelon     1.00      1.00      1.00        23
   orange      1.00      1.00      1.00        29
   papaya      1.00      1.00      1.00        19
 pigeonpeas    1.00      1.00      1.00        18
 pomegranate    1.00      1.00      1.00        17
    rice       1.00      0.75      0.86        16
 watermelon    1.00      1.00      1.00        15

 accuracy      0.99      0.99      0.99       440
 macro avg      0.99      0.99      0.99       440
 weighted avg      0.99      0.99      0.99       440
```

```
In [21]: # Cross validation score (NaiveBayes)
score = cross_val_score(NaiveBayes,features,target,cv=5)
score
```

```
Out[21]: array([0.99772727, 0.99545455, 0.99545455, 0.99545455, 0.99090909])
```

Saving trained Guassian Naive Bayes model

```
In [23]: import pickle
# Dump the trained Naive Bayes classifier with Pickle
NB_pkl_filename = '../models/NBClassifier.pkl'
# Open the file to save as pkl file
NB_Model_pkl = open(NB_pkl_filename, 'wb')
pickle.dump(NaiveBayes, NB_Model_pkl)
```

Support Vector Machine (SVM)

```
In [24]: from sklearn.svm import SVC
# data normalization with sklearn
from sklearn.preprocessing import MinMaxScaler
# fit scaler on training data
norm = MinMaxScaler().fit(Xtrain)
X_train_norm = norm.transform(Xtrain)
# transform testing data
X_test_norm = norm.transform(Xtest)
SVM = SVC(kernel='poly', degree=3, C=1)
SVM.fit(X_train_norm, Ytrain)
predicted_values = SVM.predict(X_test_norm)
x = metrics.accuracy_score(Ytest, predicted_values)
acc.append(x)
model.append('SVM')
print("SVM's Accuracy is: ", x)

print(classification_report(Ytest, predicted_values))
```

```
SVM's Accuracy is: 0.9795454545454545
precision    recall  f1-score   support

   apple      1.00      1.00      1.00        13
  banana      1.00      1.00      1.00        17
 blackgram      1.00      1.00      1.00        16
 chickpea      1.00      1.00      1.00        21
  coconut      1.00      1.00      1.00        21
   coffee      1.00      0.95      0.98        22
   cotton      0.95      1.00      0.98        20
   grapes      1.00      1.00      1.00        18
    jute       0.83      0.89      0.86        28
 kidneybeans    1.00      1.00      1.00        14
   lentil      1.00      1.00      1.00        23
```

kidneybeans	1.00	1.00	1.00	14
lentil	1.00	1.00	1.00	23
maize	1.00	0.95	0.98	21
mango	1.00	1.00	1.00	26
mothbeans	1.00	1.00	1.00	19
mungbean	1.00	1.00	1.00	24
muskmelon	1.00	1.00	1.00	23
orange	1.00	1.00	1.00	29
papaya	1.00	1.00	1.00	19
pigeonpeas	1.00	1.00	1.00	18
pomegranate	1.00	1.00	1.00	17
rice	0.80	0.75	0.77	16
watermelon	1.00	1.00	1.00	15
accuracy			0.98	440
macro avg	0.98	0.98	0.98	440
weighted avg	0.98	0.98	0.98	440

```
In [37]: # Cross validation score (SVM)
score = cross_val_score(SVM, features, target, cv=5)
score
```

```
Out[37]: array([0.97954545, 0.975      , 0.98863636, 0.98863636, 0.98181818])
```

```
In [27]: #Saving trained SVM model
```

```
In [28]: import pickle
# Dump the trained SVM classifier with Pickle
SVM_pkl_filename = '../models/SVMClassifier.pkl'
# Open the file to save as pkl file
SVM_Model_pkl = open(SVM_pkl_filename, 'wb')
pickle.dump(SVM, SVM_Model_pkl)
```

```
SVM_Model.pkl = open(SVM.pkl_filename, 'wb')
pickle.dump(SVM, SVM_Model.pkl)
# Close the pickle instances
SVM_Model.pkl.close()
```

Logistic Regression

```
In [29]: from sklearn.linear_model import LogisticRegression

LogReg = LogisticRegression(random_state=2)

LogReg.fit(Xtrain,Ytrain)

predicted_values = LogReg.predict(Xtest)

x = metrics.accuracy_score(Ytest, predicted_values)
acc.append(x)
model.append('Logistic Regression')
print("Logistic Regression's Accuracy is: ", x)

print(classification_report(Ytest,predicted_values))
```

```
Logistic Regression's Accuracy is: 0.9522727272727273
precision    recall  f1-score   support
```

apple	1.00	1.00	1.00	13
banana	1.00	1.00	1.00	17
blackgram	0.86	0.75	0.80	16
chickpea	1.00	1.00	1.00	21
coconut	1.00	1.00	1.00	21
coffee	1.00	1.00	1.00	22
cotton	0.86	0.90	0.88	20
	- - -	- - -	- - -	- -

kidneybeans	1.00	1.00	1.00	14
lentil	0.88	1.00	0.94	23
maize	0.90	0.86	0.88	21
mango	0.96	1.00	0.98	26
mothbeans	0.84	0.84	0.84	19
mungbean	1.00	0.96	0.98	24
muskmelon	1.00	1.00	1.00	23
orange	1.00	1.00	1.00	29
papaya	1.00	0.95	0.97	19
pigeonpeas	1.00	1.00	1.00	18
pomegranate	1.00	1.00	1.00	17
rice	0.85	0.69	0.76	16
watermelon	1.00	1.00	1.00	15
accuracy			0.95	440
macro avg	0.95	0.95	0.95	440
weighted avg	0.95	0.95	0.95	440

```
In [30]: # Cross validation score (Logistic Regression)
score = cross_val_score(LogReg,features,target,cv=5)
score
```

```
Out[30]: array([0.95      , 0.96590909, 0.94772727, 0.96590909, 0.94318182])
```

Saving trained Logistic Regression model

```
In [35]: import pickle
# Dump the trained Naive Bayes classifier with Pickle
LR.pkl_filename = '../models/LogisticRegression.pkl'
# Open the file to save as pkl file
LR_Model.pkl = open(LR.pkl_filename, 'wb')
pickle.dump(LogReg, LR_Model.pkl)
```

```
# Open the file to save as pkl file
SVM_Model.pkl = open(SVM_pkl_filename, 'wb')
pickle.dump(SVM, SVM_Model.pkl)
# Close the pickle instances
SVM_Model.pkl.close()
```

Logistic Regression

```
In [29]: from sklearn.linear_model import LogisticRegression

LogReg = LogisticRegression(random_state=2)

LogReg.fit(Xtrain,Ytrain)

predicted_values = LogReg.predict(Xtest)

x = metrics.accuracy_score(Ytest, predicted_values)
acc.append(x)
model.append('Logistic Regression')
print("Logistic Regression's Accuracy is: ", x)

print(classification_report(Ytest,predicted_values))
```

```
Logistic Regression's Accuracy is: 0.9522727272727273
precision    recall  f1-score   support
```

apple	1.00	1.00	1.00	13
banana	1.00	1.00	1.00	17
blackgram	0.86	0.75	0.80	16
chickpea	1.00	1.00	1.00	21
coconut	1.00	1.00	1.00	21
coffee	1.00	1.00	1.00	22
cotton	0.86	0.90	0.88	20

kidneybeans	1.00	1.00	1.00	14
lentil	0.88	1.00	0.94	23
maize	0.90	0.86	0.88	21
mango	0.96	1.00	0.98	26
mothbeans	0.84	0.84	0.84	19
mungbean	1.00	0.96	0.98	24
muskmelon	1.00	1.00	1.00	23
orange	1.00	1.00	1.00	29
papaya	1.00	0.95	0.97	19
pigeonpeas	1.00	1.00	1.00	18
pomegranate	1.00	1.00	1.00	17
rice	0.85	0.69	0.76	16
watermelon	1.00	1.00	1.00	15
accuracy			0.95	440
macro avg	0.95	0.95	0.95	440
weighted avg	0.95	0.95	0.95	440

```
In [30]: # Cross validation score (Logistic Regression)
score = cross_val_score(LogReg,features,target,cv=5)
score
```

```
Out[30]: array([0.95, 0.96590909, 0.94772727, 0.96590909, 0.94318182])
```

Saving trained Logistic Regression model

```
In [35]: import pickle
# Dump the trained Naive Bayes classifier with Pickle
LR_pkl_filename = '../models/LogisticRegression.pkl'
# Open the file to save as pkl file
LR_Model.pkl = open(LR_pkl_filename, 'wb')
pickle.dump(LogReg, LR_Model.pkl)
# Close the pickle instances
```

```

RF's Accuracy is: 0.990909090909091
precision    recall  f1-score   support

   apple      1.00      1.00      1.00        13
  banana      1.00      1.00      1.00        17
blackgram      0.94      1.00      0.97        16
 chickpea      1.00      1.00      1.00        21
   coconut      1.00      1.00      1.00        21
    coffee      1.00      1.00      1.00        22
   cotton      1.00      1.00      1.00        20
    grapes      1.00      1.00      1.00        18
     jute      0.90      1.00      0.95        28
kidneybeans      1.00      1.00      1.00        14
   lentil      1.00      1.00      1.00        23
    maize      1.00      1.00      1.00        21
    mango      1.00      1.00      1.00        26
 mothbeans      1.00      0.95      0.97        19
 mungbean      1.00      1.00      1.00        24
 muskmelon      1.00      1.00      1.00        23
    orange      1.00      1.00      1.00        29
  papaya      1.00      1.00      1.00        19
pigeonpeas      1.00      1.00      1.00        18
pomegranate      1.00      1.00      1.00        17
     rice      1.00      0.81      0.90        16
watermelon      1.00      1.00      1.00        15

 accuracy      0.99      0.99      0.99       440
 macro avg      0.99      0.99      0.99       440
weighted avg      0.99      0.99      0.99       440

```

RF's Accuracy is:

```

In [37]: # Cross validation score (Random Forest)
score = cross_val_score(RF, features, target, cv=5)
score

```

```

Out[37]: array([0.99772727, 0.99545455, 0.99772727, 0.99318182, 0.98863636])

```

Saving trained Random Forest model

```

In [38]: import pickle
# Dump the trained Naive Bayes classifier with Pickle
RF_pkl_filename = '../models/RandomForest.pkl'
# Open the file to save as pkl file
RF_Model_pkl = open(RF_pkl_filename, 'wb')
pickle.dump(RF, RF_Model_pkl)
# Close the pickle instances
RF_Model_pkl.close()

```

XGBoost

```

In [39]: import xgboost as xgb
XB = xgb.XGBClassifier()
XB.fit(Xtrain, Ytrain)

predicted_values = XB.predict(Xtest)

```



```
print("XGBoost's Accuracy is: ", x)

print(classification_report(Ytest,predicted_values))
```

[14:16:03] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.4.0/src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'multi:softprob' was changed from 'merror' to 'mlogloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

XGBoost's Accuracy is: 0.9931818181818182

	precision	recall	f1-score	support
apple	1.00	1.00	1.00	13
banana	1.00	1.00	1.00	17
blackgram	1.00	1.00	1.00	16
chickpea	1.00	1.00	1.00	21
coconut	1.00	1.00	1.00	21
coffee	0.96	1.00	0.98	22
cotton	1.00	1.00	1.00	20
grapes	1.00	1.00	1.00	18
jute	1.00	0.93	0.96	28
kidneybeans	1.00	1.00	1.00	14
lentil	0.96	1.00	0.98	23
maize	1.00	1.00	1.00	21
mango	1.00	1.00	1.00	26
mothbeans	1.00	0.95	0.97	19
mungbean	1.00	1.00	1.00	24
muskmelon	1.00	1.00	1.00	23
orange	1.00	1.00	1.00	29
papaya	1.00	1.00	1.00	19
pigeonpeas	1.00	1.00	1.00	18
pomegranate	1.00	1.00	1.00	17
rice	0.94	1.00	0.97	16
watermelon	1.00	1.00	1.00	15

In [46]:

```
# Cross validation score (XGBoost)
score = cross_val_score(XB, features, target, cv=5)
score
```

[08:54:44] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.4.0/src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'multi:softprob' was changed from 'merror' to 'mlogloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

[08:54:45] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.4.0/src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'multi:softprob' was changed from 'merror' to 'mlogloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

[08:54:46] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.4.0/src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'multi:softprob' was changed from 'merror' to 'mlogloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

[08:54:47] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.4.0/src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'multi:softprob' was changed from 'merror' to 'mlogloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

[08:54:48] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.4.0/src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'multi:softprob' was changed from 'merror' to 'mlogloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

Out[46]: array([0.99318182, 0.99318182, 0.99318182, 0.99090909, 0.99090909])


Saving trained XGBoost model

In [40]:

```
import pickle
# Dump the trained Naive Bayes classifier with Pickle
XB_pkl_filename = '../models/XGBoost.pkl'
# Open the file to save as pkl file
XB_Model_pkl = open(XB_pkl_filename, 'wb')
pickle.dump(XB, XB_Model_pkl)
```

```
sns.barplot(x = acc, y = model, palette='dark')
```

Out[41]:



In [42]: